

(EC-2 Makeup)

Q.1 Consider the following dataset with 4 records.

[4+2 = 6 Marks]

Input X	Output Y
1	$\exp(2)$
2	$\exp(4)$
3	$\exp(6.3)$
4	$\exp(9.2)$

Assume output $y = e^{(\alpha * x)}$. Using linear regression,

(a) Find the best value of α .

(b) Find the optimal total sum of square error.

Solution:

Camlin Page
Date / /

1. Solution :-

a) Find the best value of α

$$\ln(y_i) = \alpha x_i$$
$$J(\alpha) = (\alpha - 2)^2 + (2\alpha - 4)^2 + (3\alpha - 6.3)^2 + (4\alpha - 9.2)^2$$

Taking derivative of $J(\alpha)$ w.r.t α and equating to 0.

$$(\alpha - 2) + 2(2\alpha - 4) + 3(3\alpha - 6.3) + 4(4\alpha - 9.2) = 0$$

or

$$\alpha = 65.7/30$$
$$= 2.2$$

b) Find the optimal total sum of square error.

$$(\exp(2.2) - \exp(2))^2 + (\exp(4.4) - \exp(4))^2 + (\exp(6.6) - \exp(6.3))^2 + (\exp(8.8) - \exp(9.2))^2$$

Q.2 Consider inputs x_i which are real valued attributes and the outputs y_i which are real valued of the form $y_i = f(x_i) + e_i$, where $f(x_i)$ is the true function and e_i is a random variable representing laplacian noise with PDF given by

$$f(y_i/\theta) = \frac{1}{2\theta} * e^{-\frac{|y_i - \mu|}{\theta}}$$

Implementing a linear regression model of the form $h(x_i) = \sum_{i=0}^n \theta_i x_i$, and $\mu = h(x_i)$, find the maximum likelihood estimator of θ . Comment on the loss function. [4+1 Marks]

$$\begin{aligned} f(y_i|\theta) &= \frac{1}{2\theta} e^{-\frac{|y_i - \mu|}{\theta}} \\ L(\theta|y) &= \prod_{i=1}^n f(y_i|\theta) \\ &= \prod_{i=1}^n \frac{1}{2\theta} e^{-\frac{|y_i - \mu|}{\theta}} \\ \ln L(\theta|y) &= \ln \prod_{i=1}^n \frac{1}{2\theta} e^{-\frac{|y_i - \mu|}{\theta}} \\ &= \sum_{i=1}^n \ln \frac{1}{2\theta} e^{-\frac{|y_i - \mu|}{\theta}} \\ &= \sum_{i=1}^n \left[-\ln(2\theta) - \frac{|y_i - \mu|}{\theta} \right] \\ &= -n \ln 2\theta - \sum_{i=1}^n \frac{|y_i - \mu|}{\theta} \\ &= \underset{\theta}{\operatorname{argmax}} \left[-n \ln 2\theta - \sum_{i=1}^n \frac{|y_i - \mu|}{\theta} \right] \\ &= - \underset{\theta}{\operatorname{argmin}} \left[n \ln 2\theta + \sum_{i=1}^n \frac{|y_i - \mu|}{\theta} \right] \\ \frac{d}{d\theta} \ln L(\theta|y) &= \frac{d}{d\theta} \left[n \ln 2\theta + \sum_{i=1}^n \frac{|y_i - \mu|}{\theta} \right] \\ &= \frac{n}{\theta} - \frac{1}{\theta^2} \sum_{i=1}^n |y_i - \mu| = 0 \\ \Rightarrow \frac{n}{\theta} &= \frac{1}{\theta^2} \sum_{i=1}^n |y_i - \mu| \end{aligned}$$

$$\theta = \frac{1}{n} \sum_{i=1}^n |y_i - \mu|$$

Comment on Loss function: Instead of MSE, MAE is the maximum likelihood hypothesis. So MAE is appropriate for the loss function.

Q.3 Consider a result prediction system where student's efforts are encoded as percent of time a student has spent studying out of total available time.

- The input X is having just one feature representing the student's efforts having only four discrete values (25%, 50%, 75%, and 100%)
- The output Y is having 3 classes (First class, Second class, Fail)
- The priors for each class are: $P(Y = \text{First Class}) = 0.5$, $P(Y = \text{Second class}) = 0.3$, and $P(Y = \text{Fail}) = 0.2$.
- Based on the past data, the estimated the class-conditional probability $P(X|Y)$ are shown in the following table.

Student's efforts	$p(x y=\text{fail})$	$p(x y=\text{second class})$	$p(x y=\text{first class})$
25	0.7	0.4	0.1
50	0.2	0.3	0.1
75	0.1	0.2	0.3
100	0	0.1	0.7

Consider a following loss function $\iota(\hat{y}, y)$ where $\hat{y} = \text{predicted class label}$ and y is true class label:

$$\iota(\hat{y}, y) = \begin{cases} 0 & \hat{y} = y \\ 1 & \hat{y} = \text{Fail and } \hat{y} \neq y \\ 2 & \hat{y} = \text{Second class and } \hat{y} \neq y \\ 4 & \hat{y} = \text{First class and } \hat{y} \neq y \end{cases}$$

Consider modified Naïve Bayes hypothesis function:

$$\hat{Y} \leftarrow \underset{y_k}{\operatorname{argmax}} l(y, \hat{y}) P(Y = y_k) \prod_i P(X_i | Y = y_k)$$

Use this modified hypothesis function to classify each of the examples in the given table. [5 Marks]

Solution:

	$p(y x)$		
	Fail	second class	first class
25	0.35	0.12	0.02
50	0.1	0.09	0.02
75	0.05	0.06	0.06
100	0	0.03	0.14

	$L(y\text{-hat}, y) * p(y/x)$	$L(y\text{-hat}, y) * p(y/x)$	$L(y\text{-hat}, y) * p(y/x)$		
	Fail	second class	first class	Highest value	
25	0.35	0.24	0.08	0.35	fail
50	0.1	0.18	0.08	0.18	second class
75	0.05	0.12	0.24	0.24	first class

10					
0	0	0.06	0.56	0.56	first class

Q.4 If we modify the loss function of the linear regression model as follows:

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n w^{(i)} \left(h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

Where $w^{(i)}$ is the weight assigned to each training example. Derive the equation to find the value of β with this modified loss function. Suppose, we estimate the value of $w^{(i)}$ inversely proportional to the variance of the residuals, comment in **no more than 20 words** when you prefer to use this kind of modified loss function.

[3+2=5 Marks]

Solution:

4. Solution:-

Hint:- you might need to use a matrix W such that $\text{diag}(W) = [w_1, w_2, \dots, w_N]^T$

Define $Y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}$ and $X = \begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_N^T \end{bmatrix}$

Then $L(\beta) = (Y - X\beta)^T W (Y - X\beta)$.

Setting $\frac{dL(\beta)}{d\beta} = 0$, we get

$$\hat{\beta} = (X^T W X)^{-1} X^T W Y$$

Comment: Robust against outliers. Outliers will have higher variance of the residuals resulting into lower weight.

Q.5 Fit a logistic regression. Find the updated weights after the 3 iterations of modified Gradient Descent algorithm where gradient update happens after every training example using a learning rate of 0.5 and initial weights $(W_0, W_1, W_2) = (1, 1, 1)$ for the following data with the logistic regression output given by

$$\frac{1}{1 + e^{(-W_0 + W_1 X_1^2 - W_2 X_2)}}$$

Assume the results obtained after 3 iterations is the final weights. Using this construct

the confusion matrix for given below training data.

[4+2=6

Marks]

Input X1	Input X2	Output Label
2	0	0
0	2	0
0	-2	0
-2	0	0
0	1	1
0	-1	1

Solution:

Wi-LR*[Y-Pred - Y]*Xi						
X1	X2	y	Y-Pred = h(X)	W0	w1	w2
2	0	0	0.05	0.975	0.95	1
0	2	0	0.95	0.5	0.95	0.05
0	-2	0	0.27	0.365	0.95	0.32
-2	0	0	0.05			

w0	w1	w2
0.365	0.95	0.32

X1	X2	y	Y-Pred = h(X)	Y
2	0	0	0.03	0
0	2	0	0.73	1
0	-2	0	0.43	0
-2	0	0	0.03	0
0	1	1	0.66	1
0	-1	1	0.51	1

		True Class	
		Y=0	Y=1
Predicted Class	Y=0	3	0
	Y=1	1	2

Q.6 Consider the following set of training examples:

Instance	Classification	A ₁	A ₂
1	+	T	T
2	+	T	T
3	-	T	F
4	+	F	F
5	-	F	T
6	-	F	T

What is the information gain of A₂ relative to these training examples? Provide the equation for calculating the information gain as well as intermediate results. [3 Marks]

Camlin Page
 Date / /

6. Solution :-

$$\text{Entropy } E(S) = E([3+, 3-])$$

$$= -(3/6) \log_2(3/6) - (3/6) \log_2(3/6)$$

$$\boxed{\text{Entropy } E(S) = 1}$$

$$\text{Gain}(S, A_2) = E(S) - (4/6) E(T) - (2/6) E(F)$$

$$= 1 - (4/6) - (2/6) \approx 0$$

$$E(T) = E([2+, 2-]) = 1$$

$$E(F) = E([1+, 1-]) = 1$$